

Priya Sharma

Data Engineer | Pune, India | priya.sharma@example.com | github.com/priyasharma-demo

SUMMARY

Data engineer with five years of experience building batch and streaming pipelines for retail analytics. Comfortable owning a pipeline from ingestion to the dashboard, and happiest when a flaky job becomes a boring one.

EXPERIENCE

Senior Data Engineer

Northwind Retail Analytics, Pune | March 2023 - Present

- Lead a team of four engineers maintaining 60+ Airflow DAGs that feed the merchandising warehouse.
- Moved nightly sales reconciliation from a 6-hour cron script to an incremental dbt model that finishes in 40 minutes.
- Introduced data contracts between the point-of-sale team and analytics, cutting schema-break incidents from about two a month to zero in two quarters.
- Mentor two junior engineers and run the fortnightly pipeline review.

Data Engineer

Bluefin Logistics, Bengaluru | July 2020 - February 2023

- Built the Kafka ingestion layer for shipment tracking events, handling around 2 million events per day.
- Wrote the Spark jobs that compute delivery-delay forecasts used by the operations dashboard.
- Reduced the monthly cloud bill for the analytics cluster by 30% by moving cold partitions to object storage.

PROJECTS

SHELFSENSE : STOCKOUT PREDICTION

- Side project that predicts which store shelves will run empty in the next 48 hours from point-of-sale and delivery data.
- Gradient-boosted model served through a small FastAPI service, retrained weekly by a GitHub Actions workflow.
- Won second place at the 2024 Pune Retail Data Hackathon.

TRAILMIX : OPEN-SOURCE DATA QUALITY CHECKS

- A lightweight library of column-level checks (null rate, freshness, row-count drift) that runs inside Airflow tasks.
- About 400 stars on GitHub; adopted by two teams outside Northwind.

BUSROUTE : CITY BUS DELAY TRACKER

- Weekend project that scrapes public bus GPS feeds and shows average delay per route on a map.
- Built with DuckDB and Streamlit; the whole thing runs on a Raspberry Pi at home.

SKILLS

Languages: Python, SQL, Scala, Bash

Data tools: Apache Airflow, dbt, Apache Spark, Apache Kafka, DuckDB

Databases: PostgreSQL, Snowflake, BigQuery, MongoDB

Cloud: AWS (S3, Glue, EMR), Google Cloud Platform

Practices: data modelling, data contracts, CI/CD for pipelines, on-call and incident review

EDUCATION

Bachelor of Technology in Computer Science

College of Engineering Pune | 2016 - 2020 | CGPA 8.4 / 10

Final-year thesis on anomaly detection in smart-meter electricity readings.

CERTIFICATIONS

- AWS Certified Data Analytics - Specialty (2022)
- Google Cloud Professional Data Engineer (2024)

ACHIEVEMENTS AND INTERESTS

- Speaker at PyCon India 2023 with the talk "Testing data pipelines like application code".
- Volunteers as a mentor for a weekend coding club for school students.
- Plays competitive table tennis at the district level.